

LOCALIZATION IN THE COMMUTANT, STAGE 10:

the selected Dirac operator.

Tilings, kink currents, contact memory and the He-initialised RELU network
as a critical random spectral geometry

9 October 2026

Abstract

We place the bias-free He-initialised RELU multilayer perceptron with standard Gaussian input inside the framework of random finite spectral triples (Dirac ensembles), Sahasrabudhe's two-level picture of probabilistic combinatorics at small scales, and Klartag's evolving ellipsoid. The placement has one central object, the *selected Dirac operator*: the weights form a random chiral Dirac chain on the layered neuron set, the input selects a projection of the commutative neuron algebra (its gate pattern), and the network is the resolvent of the Dirac chain compressed by that projection. The partition of input space by the selected projection is the activation tiling; its walls are an isotropic random hyperplane process at every layer (level one), and its overlap kernels are the geometry in which the next layer's walls live (level two, the lift). Proved results: (i) a Crofton formula for the tiling and the folding law at He criticality (lengths preserved, chords contract like $3\pi/l$); (ii) the *kink-current formula*: the exact mean of the network is a sum over walls of the Gaussian surface density of the wall times the downstream transport, equivalently half the expected local time of the pre-activations along a Brownian path; (iii) the third-cumulant memory of every moment chain is a record of contacts: its hub index is the index of a kink and its identity leg the pinned coordinate (Price's theorem), and a frame bound shows that its compressibility is that of its non-Hadamard leg; (iv) the neuron algebra and the modular frame of the layer state are complementary, which is the Hadamard obstruction in operator-algebraic form; (v) the Euler–Stein correction of stage 9 is an average over Gaussian balls of uniformly distributed radius in the tiling, its constant is $1/(d+2)$ for a defect of Hermite contact order d , which is the order of the concentration-function coefficient at the kink that the estimator gets wrong; (vi) neuron-pair correlations are Mehler factors, i.e. imaginary modular times of the Gibbs state of the Hermite number operator, and the depth flow of the pair temperature is marginal at linear order with a $\beta^{3/2}$ term produced by the kink, which fixes the exponent of the folding law; (vii) the mean is the unique solution of a positivity-plus-Stein moment problem whose first loop equation is the kink-current formula, and the cumulant chains are its truncations without positivity. Every statement is labelled as proved, known (with a reference), or conjectural; the identities are checked by an accompanying script.

Contents

1	The placement in one page	2
2	Two levels: random walls and their lift	4
2.1	Level one: the walls of a layer are an isotropic point process	4
2.2	Crofton: how many walls a curve crosses	4

2.3	The lift at He criticality: lengths are preserved, chords contract	5
2.4	The lift as a tower of algebras	6
3	The network as a random spectral triple	6
3.1	The chiral Dirac chain	6
3.2	The network is a compressed resolvent	7
3.3	The selected Dirac operator and the induced ensemble	7
3.4	Distance on states	8
4	The mean is a kink current	8
4.1	The kink-current formula	8
4.2	What an estimator has to compute: conditional expectations on walls	10
5	Estimators are unsharp selections; their memory is a record of contacts	10
5.1	Moment chains are unsharp selections	10
5.2	Pinning and births: the hub is a kink	11
5.3	Kinks as points: codegree and the frame bound	11
5.4	Complementarity of the neuron algebra and the layer state	12
5.5	Parity: the odd sector and broken symmetry	13
6	The clock of the selection: Mehler factors, modular time and the Euler–Stein constant	13
6.1	The Euler–Stein correction probes the estimator on balls of uniform radius	13
6.2	Hermite contact order and the Euler–Stein constant	14
6.3	Mehler factors are imaginary modular times	15
7	Bootstrap: the mean as the unique solution of a positivity problem	16
8	What the placement explains, what it predicts, and what is open	17
8.1	What the placement explains	17
8.2	What it predicts (falsifiable, theoretical)	18
8.3	Open problems	18

1 The placement in one page

Fix $W_1, \dots, W_L \in \mathbb{R}^{n \times n}$ with i.i.d. $N(0, 2/n)$ entries, put $h_0 = x \sim \gamma = N(0, I_n)$, $z_l = W_l h_{l-1}$, $h_l = \text{relu}(z_l)$; the target is $u = \mathbb{E} h_L$. The placement consists of four identifications, each of which is a statement proved below.

1. The space and the metric. The layered neuron set $X = \{0, \dots, L\} \times [n]$ carries the commutative algebra $\mathcal{A} = C(X)$ and the Hilbert space $\ell^2(X)$. The weights are one self-adjoint operator $D = \sum_l (W_l \otimes |l\rangle\langle l-1| + W_l^\top \otimes |l-1\rangle\langle l|)$, odd for the grading $(-1)^l$: a random *chiral Dirac chain*. He initialisation is the Gaussian (quadratic-action) Dirac ensemble $e^{-\frac{n}{8} \text{Tr} D^2} dD$ at the coupling where the gated transfer has unit mean-square gain (Proposition 3.2).

2. The selection. An input x selects a projection $P_x \in \mathcal{A}$ (its gate pattern), and the network is the compressed resolvent $z_L(x) = [(1 - D \downarrow P_x)^{-1}]_{L0} x$ (Theorem 3.3). The *selected Dirac operator* is $P_x D P_x$, the Dirac operator of the active subnetwork. Input space is tiled by the value of the selected Dirac operator (the activation tiling), the Gaussian input induces a probability law on compressions (a Dirac ensemble *induced by the input*), and moving the input across a wall toggles one site, a rank-one contact update of the selected geometry (Proposition 3.5). Localization of the input (stage 9) is the dynamical selection: a Gaussian ball of shrinking radius at a moving centre, whose terminal selected Dirac operator is that of a Gaussian point.

3. The two levels. At every layer the walls $\{z_{l,a} = 0\}$ are the hyperplanes of n i.i.d. isotropic normals in the previous feature space (level one: a random point process on a sphere). Their overlaps —codegrees of caps, intersections of half-spaces—are arc-cosine kernels on both the input side and the neuron side (Lemma 2.1), and that level-two geometry is the ambient geometry of the next layer’s walls (the lift). Integral geometry (Crofton) counts the walls a curve crosses (Theorem 2.2); at He criticality lengths are preserved and chords contract like $3\pi/l$ (Theorem 2.4): the selected geometry folds. The same two-level structure reappears one level up: kinks are points with marks (kink density, covariance arms), and the third-cumulant memory is a sum of rank-one atoms over these points whose overlaps are codegrees in the covariance graph (Section 5).

4. The observable and its estimators. The mean is a boundary functional of the tiling: the *kink current* $u = \sum_{l,a} \mathbb{E}[\delta(z_{l,a}) |\nabla z_{l,a}|^2 \partial_{h_l} h_L e_a]$ (Theorem 4.1), the first Schwinger–Dyson equation of the network’s Gaussian field theory (Section 7). A moment chain is an *unsharp* selection (gate projections replaced by their expectations), its third cumulant is a record of contacts (hub = kink, identity leg = pinned coordinate, Theorem 5.3), and its errors are read by a clock: Gaussian balls of uniform radius in the tiling, equivalently an exponentially distributed Ornstein–Uhlenbeck time, which turns Hermite contact order d into the Euler–Stein constant $1/(d+2)$ (Theorems 6.1, 6.2).

Sahasrabudhe [3]	Klartag [4]	Dirac ensembles [8, 6, 5, 7]	He-ReLU network
Poisson points in a box	lattice points	random Dirac D	weight rows: random walls per layer
conflict graph, codegree	contacts $\langle y, Ay \rangle = 1$	spectral triple $(\mathcal{A}, \mathcal{H}, D)$	tiling; codegree = arc-cosine kernel
lift (forget embedding)	ellipsoid A_t (metric)	metric from D (Connes)	selected Dirac $P_x D P_x$; pullback metric
nibble: $\text{Var} \leq p \sum c_y^2$	rank-one contact updates	action $\text{Tr } D^2, \text{Tr } D^4$	He = quadratic action, critical gain
small codegree $\Rightarrow$ independence	volume growth	1-cut/2-cut, symmetry breaking	folding law; odd memory sector
small-ball ρ_h at scale h	stopping when held	spectral dimension, variance	concentration coefficients at the kink

Two cautions recorded by the stage-9 workflow and kept here. The gate algebras are commutative subalgebras of $L^\infty(\mathbb{R}^n, \gamma)$, so the modular theory of their states is trivial; the noncommutative geometry of the problem lives in the pair (neuron algebra, Dirac/transport) and in the pair (neuron algebra, layer state on M_n), Sections 3 and 5.4. And Klartag’s process is the dual of a localization: there the metric is the martingale and the centre is fixed, here the centre is the

martingale and the metric (precision) deterministic; the correspondence is between the contact structures, not the processes.

2 Two levels: random walls and their lift

2.1 Level one: the walls of a layer are an isotropic point process

Condition on layers $1, \dots, l-1$. The feature map $x \mapsto h_{l-1}(x) \in \mathbb{R}^n$ is then fixed, and the rows $w_{l,a}$ of W_l are n independent $N(0, \frac{2}{n}I)$ vectors; their directions $\hat{w}_{l,a}$ are n i.i.d. uniform points on S^{n-1} , a binomial process that is Poisson to first order. Each point carries a *wall* in input space, $F_{l,a} = \{x : \langle w_{l,a}, h_{l-1}(x) \rangle = 0\}$, and each input x carries a *cap*, the open hemisphere $C_l(x) = \{u \in S^{n-1} : \langle u, h_{l-1}(x) \rangle > 0\}$; neuron a is active at x iff $\hat{w}_{l,a} \in C_l(x)$. This is the incidence structure of [3, §3.3] (points in a box, balls around them, intersections) with the rôles of the two families symmetric.

Lemma 2.1 (Sign duality and the two codegrees). *Write $\theta(u, v)$ for the angle between two vectors and $J(\theta) = \sin \theta + (\pi - \theta) \cos \theta$.*

- (i) *(input side) For inputs x, y with $\theta = \theta(h_{l-1}(x), h_{l-1}(y))$, conditionally on layers $< l$ the codegree $N_l(x, y) = \#\{a : a \text{ active at } x \text{ and at } y\}$ is $\text{Bin}(n, \frac{\pi - \theta}{2\pi})$, and the number of neurons whose gate differs at x and y is $\text{Bin}(n, \frac{\theta}{\pi})$. The weighted codegree $K_l(x, y) = \langle h_l(x), h_l(y) \rangle$ has conditional mean $\frac{1}{\pi} |h_{l-1}(x)| |h_{l-1}(y)| J(\theta)$ and conditional variance $O(1/n)$ relative to its mean squared.*
- (ii) *(neuron side) At layer 1, for two rows with angle θ_{ab} , $\gamma(a \text{ and } b \text{ active}) = \frac{\pi - \theta_{ab}}{2\pi}$ and $\mathbb{E}_x[h_{1,a} h_{1,b}] = \frac{1}{2\pi} |w_a| |w_b| J(\theta_{ab})$.*
- (iii) *(Poisson covariance) If the number of points is Poissonised to $\text{Poi}(n)$, the counts of points in two caps have $\text{Cov}(N(C_l(x)), N(C_l(y))) = n \mu(C_l(x) \cap C_l(y)) = n \frac{\pi - \theta}{2\pi}$: the overlap of two caps is the covariance of their counts, and the codegree is the pairing $\tau(p_x p_y)$ of the two cap projections in $L^\infty(S^{n-1}, \mu)$.*

Proof. For a uniform direction u and two vectors at angle θ , $\mathbb{P}(\langle u, v_1 \rangle > 0, \langle u, v_2 \rangle > 0) = \frac{\pi - \theta}{2\pi}$ (the lune between two hemispheres) and $\mathbb{P}(\text{signs differ}) = \frac{\theta}{\pi}$; the rows are independent, which gives the binomial laws. The weighted codegree is a sum of n i.i.d. terms $\text{relu}(\langle w, h(x) \rangle) \text{relu}(\langle w, h(y) \rangle)$, whose mean is the degree-one arc-cosine kernel $\frac{2}{n} \frac{|h(x)| |h(y)|}{2\pi} J(\theta)$ [11]; the variance statement is the law of large numbers for n i.i.d. terms. (ii) is the same computation with $x \sim \gamma$ isotropic and the rows fixed. (iii) is the Poisson covariance identity: split the two caps into the three disjoint pieces; counts of disjoint pieces are independent and the variance of a Poisson count equals its mean. $\square$

So the post-activation covariance of a layer (the object every moment chain propagates) is the weighted codegree of the neurons' half-spaces, and the feature kernel on inputs is the weighted codegree of the inputs' caps. Neither side is privileged: the incidence $(x, a) \mapsto \mathbf{1}[\langle w_a, h(x) \rangle > 0]$ is the sign pattern of one bilinear form.

2.2 Crofton: how many walls a curve crosses

Theorem 2.2 (Crofton formula for the tiling). *Let $c : [0, 1] \rightarrow \mathbb{R}^n$ be a piecewise C^1 curve of inputs with $h_{l-1}(c(t)) \neq 0$, and let $\Lambda_{l-1}(c) = \int_0^1 \left| \frac{d}{dt} h_{l-1}(c(t)) \right| dt$ be the angular length of its image in the layer- $(l-1)$ feature sphere. Conditionally on layers $< l$, the expected number of crossings*

of layer- l walls by c is

$$\mathbb{E}[\#\{(a, t) : z_{l,a}(c(t)) = 0\} \mid \text{layers} < l] = \frac{n}{\pi} \Lambda_{l-1}(c).$$

Proof. For one row, $X_t = \langle w, \widehat{h_{l-1}(c(t))} \rangle \sqrt{n/2}$ is a centred Gaussian process with unit variance and $\text{Cov}(X_t, \dot{X}_t) = 0$ (the derivative of a unit vector is orthogonal to it), so the Kac–Rice formula gives $\mathbb{E}\#\{t : X_t = 0\} = \frac{1}{\pi} \int_0^1 \sqrt{\text{Var } \dot{X}_t} dt = \frac{1}{\pi} \int_0^1 \left| \frac{d}{dt} \widehat{h_{l-1}} \right| dt$. The zeros of X_t are the crossings of the wall $F_{l,a}$; sum over the n rows. $\square$

This is the integral-geometric face of the two-level picture: the expected number of random walls (level one) that meet a set is the set’s mean width in the geometry of the previous level. For a chord the count of *differing gates* is $n\theta/\pi$ (Lemma 2.1); for the curve it is $n\Lambda/\pi$. The two differ as soon as the image curve folds.

2.3 The lift at He criticality: lengths are preserved, chords contract

Proposition 2.3 (Unit angular gain). *Let $v \perp c$ be a tangent vector at a feature vector c , $c' = \text{relu}(Wc)$, $\dot{c}' = D(Wc)Wv$ its image, W with i.i.d. $N(0, \sigma_w^2/n)$ entries. Then $\mathbb{E}|c'|^2 = \frac{\sigma_w^2}{2}|c|^2$, $\mathbb{E}|\dot{c}'|^2 = \frac{\sigma_w^2}{2}|v|^2$, $\mathbb{E}\langle c', \dot{c}' \rangle = 0$ and $\text{Var}\langle c', \dot{c}' \rangle = O(|c|^2|v|^2/n)$. Hence the expected squared angular speed is multiplied by $1 + O(1/n)$ for every σ_w , and lengths are preserved ($\sigma_w^2 = 2$) iff the gated transfer has unit mean-square gain, $\chi_1 = \sigma_w^2 \mathbb{E}[\text{relu}'(z)^2] = 1$.*

Proof. $\langle w_a, c \rangle$ and $\langle w_a, v \rangle$ are independent Gaussians because $v \perp c$; each neuron is active with probability $\frac{1}{2}$ independently of $\langle w_a, v \rangle$. Then $\mathbb{E}|\dot{c}'|^2 = n \cdot \frac{1}{2} \cdot \frac{\sigma_w^2}{n} |v|^2$, $\mathbb{E}|c'|^2 = n \cdot \mathbb{E}\text{relu}(g)^2$ with $g \sim N(0, \sigma_w^2|c|^2/n)$, and $\langle c', \dot{c}' \rangle = \sum_a \text{relu}(\langle w_a, c \rangle) \langle w_a, v \rangle$ is a sum of n independent centred terms of variance $\frac{\sigma_w^2|c|^2}{2n} \frac{\sigma_w^2|v|^2}{n}$. $\square$

Theorem 2.4 (Folding law; infinite width). *In the limit $n \rightarrow \infty$ with L fixed (the NNGP limit [11, 12]), the angle θ_l between the features of two inputs obeys the σ_w -independent recursion $\cos \theta_{l+1} = J(\theta_l)/\pi$, and*

$$\theta_{l+1} = \theta_l - \frac{\theta_l^2}{3\pi} + O(\theta_l^3), \quad \theta_l = \frac{3\pi}{l + 3\pi/\theta_0 + O(\log l)}.$$

The angular length of every curve is preserved, $\Lambda_l = \Lambda_0$. Consequently, for two inputs at angle θ_0 , the walls of layer l crossed by the great-circle arc between them number $n\theta_0/\pi$ in expectation at every layer (in total $Ln\theta_0/\pi$), while the gates that differ between the endpoints number $n\theta_{l-1}/\pi \approx 3n/l$ (in total $\approx 3n \log L$). The selected geometry is an isometric immersion of the input sphere into the feature sphere whose chords shrink like $3\pi/l$: it folds.

Proof. The recursion is Lemma 2.1(i) with the law of large numbers, normalised by the norms (the σ_w factor cancels in the angle because relu is positively homogeneous). Expanding, $\cos \theta' - \cos \theta = (\sin \theta - \theta \cos \theta)/\pi = \theta^3/(3\pi) + O(\theta^5)$, and $\cos \theta' = 1 - \theta'^2/2 + O(\theta'^4)$ give $\theta'^2 = \theta^2 - \frac{2\theta^3}{3\pi} + O(\theta^4)$, i.e. $\theta' = \theta - \theta^2/(3\pi) + O(\theta^3)$. Then $1/\theta_{l+1} = 1/\theta_l + 1/(3\pi) + O(\theta_l)$, and summing the $O(\theta_l) = O(1/l)$ errors gives the $O(\log l)$ term. The angle map $f(\theta) = \arccos(J(\theta)/\pi)$ has $f'(0) = 1$, so the length of a curve, $\lim \sum f(\Delta\theta_i) = \sum \Delta\theta_i$, is preserved; this is the infinite-width form of Proposition 2.3. The crossing counts follow from Theorem 2.2 and Lemma 2.1(i). $\square$

Remark 2.5 (Two metrics on inputs). The *chord* metric $|\hat{h}_l(x) - \hat{h}_l(y)| = 2 \sin(\theta_l/2)$ and the *length* metric (the infimum of image lengths of curves joining x and y) separate at depth: the first collapses like $3\pi/l$, the second is the round metric of the input sphere at every depth. Section 3.4 identifies the second as the Connes distance of the selected Dirac operator. Bias-free RELU networks are thus at the edge of chaos for every σ_w in the angle variable, and He initialisation is the point at which the length variable is critical too [12, 13]; the folding exponent is computed spectrally in Proposition 6.6. The linear-in- L crossing count is the infinite-width form of the region counts of [14].

2.4 The lift as a tower of algebras

Let $\mathcal{B}_l \subset L^\infty(\mathbb{R}^n, \gamma)$ be the von Neumann algebra generated by the gate functions $\mathbf{1}[z_{k,a} > 0]$, $k \leq l$. Then $\mathcal{B}_1 \subset \mathcal{B}_2 \subset \dots \subset \mathcal{B}_L$; $\mathcal{B}_l$ is the algebra of functions constant on the tiles of the depth- l tiling, its minimal projections are the tiles, and the trace γ gives each tile its Gaussian measure. The level-two data of layer l (the kernels K_l on inputs, the post-activation covariance C_l on neurons) are the ambient metric data of the level-one point process of layer $l+1$. The network is a $\mathcal{B}_L$ -valued piecewise-linear section: on the tile with projection p_g it is the linear map J_g , so that $h_L = \sum_g p_g J_g x$. Everything that follows is about how the walls of this tower, and not its tiles, carry the mean.

3 The network as a random spectral triple

3.1 The chiral Dirac chain

Definition 3.1 (Network spectral triple). Let $X = \{0, \dots, L\} \times [n]$, $\mathcal{A} = C(X) \cong \bigoplus_{l=0}^L \mathbb{C}^n$ acting diagonally on $\mathcal{H} = \ell^2(X)$, and

$$D = D_\downarrow + D_\downarrow^\top, \quad D_\downarrow = \sum_{l=1}^L W_l \otimes |l\rangle\langle l-1|,$$

with grading $\Gamma = \bigoplus_l (-1)^l$ and real structure $\mathcal{J} = \text{complex conjugation}$. Then $D\Gamma = -\Gamma D$, $\mathcal{J}D = D\mathcal{J}$, $\mathcal{J}\Gamma = \Gamma\mathcal{J}$, $\mathcal{J}^2 = 1$, and the order-zero and first-order conditions hold trivially because $\mathcal{A}$ is commutative: $(\mathcal{A}, \mathcal{H}, D, \Gamma, \mathcal{J})$ is a real even finite spectral triple of KO-dimension 0. Its commutators $[D, f]$, $f \in \mathcal{A}$, have entries $W_l(a, b)(f_{l-1,b} - f_{l,a})$: D is the weighted adjacency of the layered graph, and the Connes distance of the triple is the metric on neurons in which strong weights are short edges.

Proposition 3.2 (He initialisation as a critical Gaussian Dirac ensemble). *If the entries of every W_l are i.i.d. $N(0, \sigma_w^2/n)$, the law of D on the space of chiral chain operators is $Z^{-1} \exp(-\frac{n}{4\sigma_w^2} \text{Tr } D^2) dD$, the quadratic Barrett–Glaser action [8, 6] restricted to this kinematic space. For a fixed vector z and a gate projection P selecting the positive coordinates of z , $\mathbb{E}|W Pz|^2 = \sigma_w^2 |Pz|^2$ and $\mathbb{E}|Pz|^2 = \frac{1}{2}|z|^2$ for symmetric z ; the gated transfer has unit mean-square gain exactly at $\sigma_w^2 = 2$, and $\mathbb{E} \log(|z_{l+1}|^2/|z_l|^2) = \log(\sigma_w^2/2) + O(1/n)$. He initialisation is the critical coupling of the selected (gated) chain.*

Proof. $\text{Tr } D^2 = 2 \sum_l \text{Tr } W_l^\top W_l$ and the Gaussian density of the entries is $\prod \exp(-\frac{n}{2\sigma_w^2} W_{ij}^2)$. For fixed $y = Pz$, Wy has i.i.d. $N(0, \sigma_w^2 |y|^2/n)$ coordinates. The logarithmic statement is the law of $\log \chi_n^2/n$, whose mean is $-O(1/n)$. $\square$

3.2 The network is a compressed resolvent

Theorem 3.3 (Gated resolvent). *For an input x let $P_x \in \mathcal{A}$ be the projection that is the identity on layer 0 and the gate projection $\mathbf{1}_{(0,\infty)}(z_l(x))$ on layer $l \geq 1$, and $\psi = \bigoplus_l z_l$ the field of pre-activations ($z_0 = x$). Then ψ is the unique solution of the self-consistent equation*

$$\psi = x \oplus 0 + D_{\downarrow} P_{\psi} \psi, \quad \text{i.e.} \quad \psi = (1 - D_{\downarrow} P_{\psi})^{-1}(x \oplus 0),$$

where $P_{\psi} = \mathbf{1}_{(0,\infty)}(\psi)$ is a spectral projection of the field in the commutative algebra. Since $D_{\downarrow} P$ is nilpotent, $(1 - D_{\downarrow} P)^{-1} = \sum_{k=0}^L (D_{\downarrow} P)^k$, and

$$z_L(x) = [(1 - D_{\downarrow} P_x)^{-1}]_{L0} x = W_L P_{L-1} W_{L-1} \cdots P_1 W_1 x,$$

$$[z_L(x)]_k = \sum_{a_0 \rightarrow a_1 \rightarrow \cdots \rightarrow a_L = k} \prod_{l=1}^L W_l(a_l, a_{l-1}) \prod_{l=1}^{L-1} \mathbf{1}[z_{l,a_l}(x) > 0] x_{a_0}.$$

Proof. The equation is block lower triangular: its layer- l component reads $\psi_l = W_l P_{l-1} \psi_{l-1}$ with $P_{l-1} \psi_{l-1} = \text{relu}(\psi_{l-1})$, which is the forward pass; uniqueness follows by induction on l . The Neumann series terminates because $(D_{\downarrow} P)^{L+1} = 0$, and the $(L, 0)$ block of $(D_{\downarrow} P)^L$ is the displayed product, whose entries are the path sums. $\square$

The network therefore computes with one fixed random operator D and one input-dependent projection P_x of the commutative algebra. Everything nonlinear is in the selection $x \mapsto P_x$; given P_x the network is the resolvent of a compressed Dirac operator. The path expansion is the sum over walks of the counting problems of stage 9's reading (exterior-algebra path counts): the gates are input-dependent vertex weights.

3.3 The selected Dirac operator and the induced ensemble

Definition 3.4 (Selection). The *selected Dirac operator* of x is $D_x = P_x D P_x$, the Dirac operator of the active subnetwork, and the *selected metric* is the pullback $g_x = J(x)^{\top} J(x)$ of the output metric, $J(x) = \partial h_L / \partial x$. The activation tiling is the partition of $\mathbb{R}^n$ by the value of P_x ; the input law induces on the finite set of compressions $\{P_g D P_g\}$ the probability $\gamma(\text{tile } g)$. We call this law the Dirac ensemble *induced by the input*: a random finite geometry drawn from the $2^{n(L-1)}$ compressions of one random Dirac operator.

Proposition 3.5 (Walls are rank-one contacts). *Let x cross the wall $F_{l,a}$ transversally from $z_{l,a} < 0$ to $z_{l,a} > 0$, all other gates fixed. Then P_x changes by the minimal projection $e_{(l,a)}$, the selected Dirac operator changes by $e_{(l,a)} D P + P D e_{(l,a)} + e_{(l,a)} D e_{(l,a)}$ (a perturbation of rank at most two supported on the edges at the toggled site), the Jacobian changes by the rank-one operator*

$$\Delta J = (\partial_{h_l} h_L) e_a \nabla z_{l,a}^{\top},$$

and the selected metric changes by $\Delta g = \nabla z_{l,a} v^{\top} J + J^{\top} v \nabla z_{l,a}^{\top} + |v|^2 \nabla z_{l,a} \nabla z_{l,a}^{\top}$ with $v = (\partial_{h_l} h_L) e_a$, a rank-two update whose new direction is the normal of the wall.

Proof. Only the gate of (l, a) changes, so $J = \partial_{h_l} h_L P_l \partial_x z_l$ changes by $\partial_{h_l} h_L e_a e_a^{\top} \partial_x z_l$, and $e_a^{\top} \partial_x z_l = \nabla z_{l,a}^{\top}$; downstream gates are continuous across a generic wall. The other two statements are direct expansions. $\square$

This is the precise sense in which the tiling plays the rôle of Klartag’s lattice. In [4] the ellipsoid matrix A_t moves until a lattice point y reaches its boundary, and the contact imposes the rank-one constraint $\text{Tr}(Ayy^\top) = 1$. Here the selected geometry is constant inside a tile and receives a rank-one (Jacobian) or rank-two (metric) update at every wall a path crosses; along any path x_t the selected operators D_{x_t} form a Hamiltonian on a dynamically evolving graph whose edges switch at wall crossings. The localization processes of stage 9 are exactly such paths with a Gaussian ball attached: on Eldan’s path the ball $N(m_t, (1+t)^{-1}I)$ shrinks around a martingale centre m_t until it lies in one tile, and the terminal selected operator is D_X , $X \sim \gamma$. The process is dual to Klartag’s (there the metric is the martingale and the centre fixed); what transfers is the contact structure, and Section 6 shows that the estimator errors are read off exactly from it.

3.4 Distance on states

Proposition 3.6 (Connes–Kantorovich distance of the selected geometry). *Let*

$$d_J(x, y) = \inf_c \int_0^1 |J(c(t)) \dot{c}(t)| dt$$

over piecewise C^1 curves c from x to y (the length of the image curve; a pseudometric because J may have a kernel), and for probability laws ν, ν' with first moments let

$$d_D(\nu, \nu') = \sup\{|\nu(f) - \nu'(f)| : f \text{ 1-Lipschitz for } d_J\}.$$

Then $d_D = W_1^{d_J}$ (Kantorovich–Rubinstein), every output coordinate is 1-Lipschitz for d_J , and hence $|u(\nu)_k - u(\nu')_k| \leq W_1^{d_J}(\nu, \nu')$ for the exact means of any two input laws. For a smooth nondegenerate metric d_D is Connes’ spectral distance with the Dirac operator of g [19, 18]; here it is its length-space extension to the piecewise-flat selected metric. In the infinite-width limit d_J restricted to the input sphere is the round metric at every depth (Theorem 2.4), while the chord metric of the features collapses: the distance on states is preserved by the selection and the distance between features is not.

Proof. Kantorovich–Rubinstein duality on the (pseudo)metric space $(\mathbb{R}^n, d_J)$; $|h_{L,k}(x) - h_{L,k}(y)| \leq |h_L(x) - h_L(y)| \leq \int |J\dot{c}|$ for every curve. The last statement is Theorem 2.4. $\square$

4 The mean is a kink current

4.1 The kink-current formula

Write $\partial_{h_l} h_L(x)$ for the Jacobian of the output with respect to the post-activations of layer l ($= I$ for $l = L$, and $D_L W_L D_{L-1} \cdots W_{l+1}$ otherwise, with gates at x), and $f_{l,a}$ for the density of $z_{l,a}(x)$, $x \sim \gamma$.

Call a tile *dead* if some hidden layer is entirely inactive on it ($h_l \equiv 0$ there); on a dead tile the network vanishes identically and several walls coincide along its faces. The Gaussian measure of the dead tiles is $\mathbb{P}(\exists l < L : h_l(x) = 0)$, an orthant probability that is exponentially small in n ($1.6 \cdot 10^{-1}$ at $n = 6$, $8 \cdot 10^{-6}$ at $n = 16$, $L = 3$; of order 2^{-n} in general and below 10^{-300} at $n = 1024$).

Theorem 4.1 (Kink current). *For almost every weight realisation, up to a remainder carried by the faces of the dead tiles,*

$$\begin{aligned} u = \mathbb{E} h_L(x) &= \sum_{l=1}^L \sum_{a=1}^n \mathbb{E} \left[\delta(z_{l,a}(x)) |\nabla z_{l,a}(x)|^2 \partial_{h_l} h_L(x) e_a \right] \\ &= \sum_{l,a} f_{l,a}(0) \mathbb{E} \left[|\nabla z_{l,a}|^2 \partial_{h_l} h_L e_a \mid z_{l,a} = 0 \right]. \end{aligned} \quad (1)$$

Each term is the Gaussian surface integral of the wall $F_{l,a}$, weighted by the squared gradient of the pre-activation that defines it and by the transport of a unit kick at (l, a) to the output.

Proof. $F = h_L$ is piecewise linear and positively homogeneous, so Euler's relation $F = x \cdot \nabla F$ holds off the walls, and $\nabla F = J^\top$ is piecewise constant with jumps across the walls. Gaussian integration by parts for the bounded variation field J gives $\mathbb{E} F = \mathbb{E}[x \cdot \nabla F] = \mathbb{E}[\text{div } \nabla F] = \mathbb{E}[\Delta F]$, where ΔF is the distributional Laplacian, a measure carried by the walls. Write $J = D_L W_L D_{L-1} \cdots D_1 W_1$ and differentiate the gates: $\partial_i D_l = \sum_a \delta(z_{l,a}) \partial_i z_{l,a} e_a e_a^\top$. By the product rule

$$\sum_i (\partial_i J) \cdot i = \sum_{l,a} \delta(z_{l,a}) (\partial_{h_l} h_L) e_a \sum_i \partial_i z_{l,a} (e_a^\top W_l D_{l-1} \cdots W_1)_i = \sum_{l,a} \delta(z_{l,a}) |\nabla z_{l,a}|^2 (\partial_{h_l} h_L) e_a,$$

because $e_a^\top W_l D_{l-1} \cdots W_1 = \nabla z_{l,a}^\top$. Distinct walls meet in sets of codimension two for almost every weight realisation, the downstream gates are continuous across a generic point of $F_{l,a}$, and $\{z_{l,a} = 0\}$ is a finite union of pieces of hyperplanes (on each live tile $z_{l,a}$ is a nonzero linear form almost surely), so the products of distributions are well defined and $\mathbb{E}[\delta(z)G] = f(0) \mathbb{E}[G \mid z = 0]$ with a continuous density f . Taking expectations gives (1). The only exception is a face of a dead tile: there a downstream pre-activation vanishes identically on one side, several gates jump at once, the product rule must be applied to the simultaneous jump, and $z_{l,a}$ has an atom at 0; these faces carry the remainder, which is bounded by the Gaussian surface measure of the dead tiles' boundaries times the size of the transport. (The identity check confirms both halves: at $n = 16$ the two sides agree to the Monte Carlo noise, 1.5%; at $n = 6$, where dead tiles carry 16% of the mass, they differ by 25%.) $\square$

Example 4.2. For $L = 1$: $f_{1,a}(0) = 1/(\sqrt{2\pi}|w_a|)$, $|\nabla z_{1,a}|^2 = |w_a|^2$, so $u_a = |w_a|/\sqrt{2\pi}$, the mean of $\text{relu}(N(0, |w_a|^2))$. For $L = 2$ the layer-1 terms give $D_2 W_2$ applied to the layer-1 kink masses on the walls, and the layer-2 terms give the kink masses of the curved walls $\{w_{2,a} \cdot h_1(x) = 0\}$.

Theorem 4.3 (Tanaka form: the mean is transported local time). *Let B be a standard Brownian motion in $\mathbb{R}^n$ started at 0 and $L_t^{(l,a)}$ the semimartingale local time at 0 of $z_{l,a}(B_t)$. Then $h_L(B_t)$ is a semimartingale with*

$$\begin{aligned} dh_L(B_t) &= J(B_t) dB_t + \frac{1}{2} \sum_{l,a} (\partial_{h_l} h_L)(B_t) e_a dL_t^{(l,a)}, \\ u = \mathbb{E} h_L(B_1) &= \frac{1}{2} \sum_{l,a} \mathbb{E} \int_0^1 (\partial_{h_l} h_L)(B_t) e_a dL_t^{(l,a)}. \end{aligned}$$

In particular the kink mass $m_{l,a} = \mathbb{E}[\delta(z_{l,a})|\nabla z_{l,a}|^2]$ is half the expected local time that $z_{l,a}$ accumulates at its kink along a unit-time Brownian path from the origin.

Proof. h_L is a difference of convex functions on each line, piecewise linear, and the Itô–Tanaka (Meyer) formula applies: $dh_L(B) = \nabla h_L \cdot dB + \frac{1}{2} \Delta h_L(dt)$ with the Laplacian measure computed in the proof of Theorem 4.1; the occupation-time formula converts $\delta(z_{l,a}) |\nabla z_{l,a}|^2 dt$ into $dL^{(l,a)}$. Homogeneity gives $\mathbb{E}[\delta(z(B_t)) |\nabla z(B_t)|^2 G(B_t)] = t^{-1/2} \mathbb{E}[\delta(z) |\nabla z|^2 G]$ for degree-zero G , and $\frac{1}{2} \int_0^1 t^{-1/2} dt = 1$, consistent with (1). $\square$

Corollary 4.4 (Mean is created only at kinks). *For every neuron, $\mathbb{E}h_{l,a} = m_{l,a} + \mathbb{E}[\mathbf{1}[z_{l,a} > 0] \Delta z_{l,a}]$, where $\Delta z_{l,a}$ is the wall measure of the lower layers transported to $z_{l,a}$. A layer adds mean exactly through its kink masses and transports the lower layers’ kink currents through its gates; a network without kinks has zero mean ($\mathbb{E}z_l = \mathbb{E}\Delta z_l$ and Δ of a linear map vanishes).*

Proof. Apply the computation of Theorem 4.1 to the single coordinate $h_{l,a}$: $\Delta h_{l,a} = \delta(z_{l,a}) |\nabla z_{l,a}|^2 + \mathbf{1}[z_{l,a} > 0] \Delta z_{l,a}$. $\square$

4.2 What an estimator has to compute: conditional expectations on walls

Formula (1) reduces the mean to two kinds of quantities per wall: the density $f_{l,a}(0)$ of the pre-activation at its kink, and a conditional expectation *on the wall*, a measure-zero event. The first is the $h \rightarrow 0$ limit of the concentration function $Q_{l,a}(h) = \mathbb{P}(|z_{l,a}| \leq h) = 2f_{l,a}(0)h + \frac{1}{3}f_{l,a}''(0)h^3 + \dots$ (the pre-activation law is absolutely continuous with a continuous density, small-ball exponent one, no arithmetic structure: the gates are functions of a continuous Gaussian). On the official network 0 the Gaussian surrogate $\varphi(\alpha)/\sigma$ of $f_{l,a}(0)$ is right to $1-3 \cdot 10^{-3}$ at every layer when the mean and variance are right (stage-9 workflow, $2.6 \cdot 10^5$ inputs), so the densities are not where estimators fail. The second kind is a conditional expectation relative to an event of probability zero, the extreme case of the rare-event-relative conditioning of [3, §4] and [10]: the error tolerated in $\mathbb{E}[\cdot | z_{l,a} = 0]$ is relative to the kink mass, not absolute. For a Gaussian state, conditioning on $z_{l,a} = 0$ is a rank-one *pinning* (a directional localization along $\nabla z_{l,a}$ taken to its end), and Section 5 shows that the third cumulant of every moment chain is the record of exactly these pinnings.

5 Estimators are unsharp selections; their memory is a record of contacts

5.1 Moment chains are unsharp selections

A moment chain does not know P_x ; it knows the Gaussian state $N(m_l, S_l)$ of each layer and replaces the gate projection by its expectation, the diagonal effect $\Phi_l = \text{diag } \Phi(\alpha_{l,a}) \in \mathcal{A}$, $0 \leq \Phi_l \leq 1$. Its transport is $T_{l \leftarrow l'} = W_l \Phi_{l-1} W_{l-1} \cdots \Phi_{l'} W_{l'+1}$, the resolvent of the Dirac chain compressed by an effect instead of a projection.

Proposition 5.1 (Unsharp selection). *With the path notation of Theorem 3.3,*

$$\mathbb{E}_x[\partial_{h_{l'}} z_l(x)] - T_{l \leftarrow l'} = \sum_{\text{paths}} \prod W \left(\mathbb{E}[\prod_k \mathbf{1}[z_{k,a_k} > 0]] - \prod_k \Phi_{k,a_k} \right),$$

the sum over paths of the connected gate correlations along the path. The chain is exact for the expected transport iff the gates on every path are uncorrelated; its correction terms are correlations between the projection and the field that selects it (measurement back-action).

Proof. Take expectations in the path expansion term by term. $\square$

Only the on-wall transports of (1) enter the mean, so what matters is the back-action *conditional on a kink*. That is where the third cumulant comes from.

5.2 Pinning and births: the hub is a kink

Lemma 5.2 (Pinning). *If $z \sim N(m, S)$, then z conditioned on $z_a = 0$ is $N(m - S e_a m_a / S_{aa}, S - S e_a e_a^\top S / S_{aa})$: a rank-one update of the state along the covariance column of the pinned neuron. Downstream gate expectations shift to first order by $-\varphi(\alpha_b) \rho_{ab} \alpha_a$, $\rho_{ab} = S_{ab} / \sqrt{S_{aa} S_{bb}}$.*

Proof. Gaussian conditioning; $\Phi((m_b - S_{ab} m_a / S_{aa}) / \sigma_b) - \Phi(\alpha_b) = -\varphi(\alpha_b) \rho_{ab} \alpha_a + O(\rho^2)$. $\square$

Theorem 5.3 (Births are contacts). *Let $z \sim N(m, S)$ with small off-diagonal covariances and $h = \text{relu}(z)$. For distinct i, j, k ,*

$$\kappa_3(h_i, h_j, h_k) = \sum_{c \in \{i, j, k\}} \Phi(\alpha_a) S_{ac} \frac{\varphi(\alpha_c)}{\sigma_c} S_{cb} \Phi(\alpha_b) \Big|_{\{a, b\} = \{i, j, k\} \setminus \{c\}} + O(S_{\text{off}}^3),$$

the sum over which of the three neurons is the hub. The hub contributes through $\mathbb{E}[\text{relu}''(z_c)] = \varphi(\alpha_c) / \sigma_c = f_c(0)$, its kink density, and the two arms are the covariances to the hub weighted by the gates. In CP form, $\kappa_3 = \text{Sym} \sum_c (\Phi \circ S_c) \otimes 3e_c \otimes (f_c(0) S_c \circ \Phi)$ to second order in S_{off} , which is the star block of the organisers' chain $(\Phi \circ S^{\text{off}}, 3I, (w_2 \circ S^{\text{off}} \circ \Phi)^\top)$, $w_2 = f(0)$; its residual block makes the birth exact beyond second order.

Proof. Joint cumulants of functions of a Gaussian vector expand over connected graphs on the index set with edges weighted by off-diagonal covariances and vertex factors $\mathbb{E} f^{(\text{deg})}(z_v)$ (Price's theorem, or Wick's theorem applied to the Hermite expansions). On three vertices the connected graphs with two edges are the three paths, centred at the hub, with vertex factors $\mathbb{E} \text{relu}'(z_a) = \Phi(\alpha_a)$, $\mathbb{E} \text{relu}''(z_c) = \varphi(\alpha_c) / \sigma_c$, $\mathbb{E} \text{relu}'(z_b) = \Phi(\alpha_b)$; the triangle is $O(S_{\text{off}}^3)$. Writing the hub as the middle leg e_c gives the CP form, and Sym distributes the three positions (factor 3). $\square$

So the identity leg that makes the third-cumulant memory expensive is not an accident of a factorisation: it is the label of the kink at which the cumulant was born. Every birth block is n rank-one atoms, one per neuron c of the birth layer, with intensity $f_c(0)$ and arms S_c ; transported, the atom of hub c born at l' and read at l is $u_c z_c^\top$ with $u_c = (T A e_c) \circ (T e_c)$ and $z_c = T C e_c$, $T = T_{l \leftarrow l'}$. Lemma 5.2 says what the atom is: the first-order response of the downstream state to pinning neuron c at its kink. Measured on network 0 (stage-9 workflow): the atom energies are proportional to $f_c(0)^2$ (log-log correlation 0.97–0.998 at ages 1, 3, 6), and the atoms are mutually orthogonal to 10^{-3} ($\|\sum_c P^{(c)}\|^2 / \sum_c \|P^{(c)}\|^2 = 0.998\text{--}1.004$).

5.3 Kinks as points: codegree and the frame bound

The second appearance of the two levels: the kinks of a layer are a marked point process on the neuron index (marks: density $f_c(0)$, arms S_c), and the memory is the sum of their atoms. Two kinks overlap through the neurons their arms share, $\langle S_c \circ \Phi, S_{c'} \circ \Phi \rangle = (S \Phi^2 S)_{cc'}$: a codegree in the covariance graph. At He initialisation the off-diagonal covariances are $O(n^{-1/2})$, so codegrees are $O(n^{-1/2})$ of the degrees, and the Gram matrix of the atoms, a Hadamard product of two such codegree matrices, has off-diagonal entries $O(n^{-1})$ of the diagonal. This is exactly the small-codegree regime of [3, §3.3–3.4]. There it is what makes the nibble work (the conditional variance of a surviving degree is at most $p \sum_y c_y^2$); here the same orthogonality makes the memory incompressible.

Theorem 5.4 (Frame bound). *Let $P = UZ^\top = \sum_{c=1}^N u_c z_c^\top$ with $U, Z \in \mathbb{R}^{n \times N}$, and let $\hat{P} = U M Z^\top$ with $\text{rank } M \leq k$ be any compression in the hub index (keeping k hubs, reweighting, or mixing hubs through any k -dimensional subspace). Then*

$$\|P - \hat{P}\|_F^2 \geq \lambda_{\min}(U^\top U) \|(I - M)Z^\top\|_F^2 \geq \lambda_{\min}(U^\top U) \lambda_{\min}(Z^\top Z) (N - k).$$

Proof. $\|U(I - M)Z^\top\|_F^2 = \text{tr}(Z(I - M)^\top G_u (I - M)Z^\top) \geq \lambda_{\min}(G_u) \|(I - M)Z^\top\|_F^2$ with $G_u = U^\top U$, and $\|(I - M)Z^\top\|_F^2 = \text{tr}((I - M)G_z(I - M)^\top) \geq \lambda_{\min}(G_z) \|I - M\|_F^2 \geq \lambda_{\min}(G_z)(N - k)$ by Eckart–Young. $\square$

Corollary 5.5 (The memory is as compressible as its transport leg, no more). *If the Hadamard leg is well conditioned, $U^\top U \approx cI$ (orthogonal hub atoms of comparable size), every hub-index compression loses at least c times the energy of the transport leg $Z = TC$ outside the kept subspace. The cheapest compression is therefore the low-rank truncation of TC , and its error is that truncation’s error. On network 0 the Hadamard bulk is near-isotropic while TC keeps 91/96/99% of the slice energy at rank 64/128/256 (stage-9 workflow), against the $1 - 10^{-4}$ of energy that a 10^{-2} relative slice accuracy requires; this is the measured failure of every tier in stage 9, now with its reason.*

5.4 Complementarity of the neuron algebra and the layer state

The gate algebras are commutative and the modular theory of their states is trivial. The genuinely noncommutative pair is the neuron algebra $A = \text{diag}(\mathbb{R}^n) \subset M_n$, where the nonlinearity acts (gates, Hadamard products, functional calculus), and the layer state $\rho_l = M_l / \text{Tr } M_l$ on M_n , $M_l = \mathbb{E}[z_l z_l^\top]$, faithful and non-tracial, whose modular group $\sigma_s(Y) = \rho_l^{is} Y \rho_l^{-is}$ is the frame in which the transport acts.

Proposition 5.6 (Takesaki for the neuron algebra). *The only conditional expectation of M_n onto A is the pinching $\Delta(Y) = \text{diag}(Y)$, and it preserves ρ iff ρ is diagonal; equivalently (Takesaki [17]) A is σ^ρ -invariant iff $\rho \in A$. The independent-neuron (mean-field) coarse-graining is state-preserving only for uncorrelated layers.*

Proof. For a conditional expectation E onto A and $i \neq j$, bimodularity gives $E(e_{ij}) = e_{ii}E(e_{ij})e_{jj} = 0$ since $E(e_{ij})$ is diagonal; so $E = \Delta$. $\text{Tr}(\rho \Delta Y) = \text{Tr}(\Delta(\rho)Y)$, which equals $\text{Tr}(\rho Y)$ for all Y iff $\rho = \Delta(\rho)$. $\square$

Theorem 5.7 (Complementarity). *Let B be the masa of the eigenprojections of ρ_l (simple spectrum), U its eigenbasis, and $C_{ai} = |U_{ai}|^2 - \frac{1}{n}$. With $\tau = \frac{1}{n} \text{Tr}$, $E_A = \Delta$ and E_B the pinching in the eigenbasis, $\|E_B E_A - J\|_{L^2(\tau) \rightarrow L^2(\tau)} = \|C\|$, $J(Y) = \tau(Y)1$. Every ρ_l -preserving conditional expectation $E_{\mathcal{N}}$ onto a subalgebra $\mathcal{N} \subseteq B$ (any spectral coarse-graining of the layer state) factors through E_B , and for every $Y \in A$ with $\tau(Y) = 0$,*

$$\|E_B Y\|_2 \leq \|C\| \|Y\|_2.$$

If the eigenvectors are Haar-distributed on the orthogonal group, $\|C\| \approx 2\sqrt{2/n}$.

Proof. The first identity is the alternating-expectation theorem for two partitions of unity into minimal projections (verified in the stage-9 workflow; it follows from writing $E_A = VV^*$, $E_B = WW^*$ with isometries onto the partitions with $V^*W = C + \frac{1}{n}\mathbf{1}\mathbf{1}^\top$, so that $E_B E_A - J = WC^\top V^*$ has norm $\|C\|$). Subalgebras of B are invariant under σ^ρ because $\rho \in B$, so ρ -preserving expectations onto them exist; each is the identity on B -valued data after E_B (for $\mathcal{N}$ spanned by blocks

Q_k of eigenprojections, $E_{\mathcal{N}}^{\rho}(Y) = \sum_k \text{Tr}(\rho Q_k Y) / \text{Tr}(\rho Q_k) Q_k$ and $\text{Tr}(\rho Q_k Y) = \text{Tr}(\rho Q_k E_B(Y))$. For $Y \in A$ with $\tau(Y) = 0$, $E_B Y = (E_B E_A - J)Y$. For Haar orthogonal eigenvectors $n|U_{ai}|^2$ is asymptotically χ_1^2 , so nC has centred entries of variance 2 and the norm of an $n \times n$ such matrix is $\approx 2\sqrt{2n}$. $\square$

At He initialisation the neuron algebra and the modular frame of the layer state are nearly mutually unbiased: a compression that respects the state (keeps eigen-directions of the covariance) sees a neuron-diagonal observable (a gate, a Hadamard product, a slice’s hub diagonal) only to $O(n^{-1/2})$ ($2\sqrt{2/n} \approx 0.09$ at $n = 1024$), and a compression that respects the neurons does not preserve the state. This is the Hadamard obstruction of stage 9 in operator-algebraic form, and it is why the spectral birth truncations of stage 9 saturate (rank 16 already keeps the spectral content and still loses a factor three): spectral information and readout information live in complementary masas.

5.5 Parity: the odd sector and broken symmetry

Proposition 5.8 (Odd sector). *Fix layers $1, \dots, l-1$. The readout at layer l (before its gates) of any third-cumulant atom born at $l' < l$ is a homogeneous cubic polynomial in W_l , hence odd under $W_l \mapsto -W_l$, and its conditional mean over the Gaussian law of W_l vanishes. Any surrogate that depends on W_l only through even statistics (spectra, $W_l^\top W_l$, the NNGP recursion, free-probability models) assigns it zero.*

Proof. Each of the three legs TAe_c , Te_c , TCe_c contains W_l exactly once as the last factor, and no gate follows it. $\square$

In the language of [5]: the Gaussian weight ensemble has the symmetry $W_l \mapsto -W_l$ for every l , the ensemble-averaged (free, NNGP) description is the symmetric phase, and each realisation carries an odd component of self-averaging size and realisation-specific sign, as the broken-symmetry equilibrium measures of the $(1, 0)$ ensemble do while the mean density of states stays symmetric. The analogy is exact in its consequence (ensemble surrogates cannot see the odd sector; the stage-9 workflow measured that its effect on the final mean is still half even, through odd readout coefficients, so surrogates capture at most a factor two of a memory worth a factor one hundred) and inexact in its mechanism: here the breaking is explicit at relative order $n^{-1/2}$ (a random field), not spontaneous at order one. The collective scalars that the chain must track exactly—the radial fourth-cumulant scalar, worth a factor 16, the trace and the collective eigenvalue of the layer covariance—play the rôle of the multi-trace couplings $(\text{tr } H^2)^2$, $\text{tr } H \text{tr } H^3$ of the Barrett–Glaser actions, whose values decide the phase there.

6 The clock of the selection: Mehler factors, modular time and the Euler–Stein constant

6.1 The Euler–Stein correction probes the estimator on balls of uniform radius

Stage 9 (Corollary *Eldan’s path*) gives, for a chain E exact on point masses and scale covariant, with $e(y) = E(N(y, I))$ and $\delta(y) = e(y) - y \cdot \nabla e(y) - \Delta e(y)$, $u - e(0) = -\frac{1}{2} \int_0^\infty (1 + \tau)^{-3/2} \mathbb{E}_{y \sim N(0, \tau I)} \delta(y) d\tau$.

Theorem 6.1 (Uniform radius). *With $r = (1 + \tau)^{-1/2} \in (0, 1]$,*

$$u - e(0) = - \int_0^1 \mathbb{E} \left[\delta(M_r/r) \right] dr, \quad M_r \sim N(0, (1 - r^2)I).$$

The localized laws on Eldan's path are the Gaussian balls $N(M_r, r^2 I)$ of radius r about a Gaussian centre, and under the defect weight their radius is uniformly distributed on $[0, 1]$; equivalently the Ornstein–Uhlenbeck time $t = -\log r$ of the localization is a standard exponential variable.

Proof. $d\tau = -2r^{-3}dr$, so $\frac{1}{2}(1 + \tau)^{-3/2}d\tau = \frac{1}{2}r^3 \cdot 2r^{-3}|dr| = dr$. On Eldan's path the localized law at time τ is $N(m, sI)$ with $s = (1 + \tau)^{-1} = r^2$ and $y = m/\sqrt{s} \sim N(0, \tau I)$, i.e. $m \sim N(0, r^2 \tau I) = N(0, (1 - r^2)I)$. Finally $\mathbb{P}(-\log r > t) = \mathbb{P}(r < e^{-t}) = e^{-t}$. $\square$

This is the precise form of “an ellipsoid in a tiling”: the correction is the average, over Gaussian balls of every radius $r \in (0, 1)$ with equal weight, of the estimator's defect on that ball. A ball of radius r meets a wall through its centre's neighbourhood with probability proportional to r (Crofton), and the order at which the estimator mis-accounts the contact decides how fast its defect vanishes as the ball shrinks.

6.2 Hermite contact order and the Euler–Stein constant

Theorem 6.2 (Contact order law). *Suppose the defect has the kink profile of order d at a wall through the origin, $\delta(y) = c(\varphi \text{He}_d)(\langle v, y \rangle)$ with $|v| = 1$ (the form of every layer-one wall of a bias-free network). Then $\mathbb{E}_{y \sim N(0, \tau I)} \delta(y) = r^{d+1} \delta(0)$, and the merge constant defined by $u - e(0) = -a \delta(0)$ is*

$$a = \int_0^1 r^{d+1} dr = \frac{1}{d+2} = \mathbb{E}[e^{-(d+1)t}], \quad t \sim \text{Exp}(1).$$

For odd d the defect vanishes at the origin and contributes nothing.

Proof. $\varphi \text{He}_d = (-1)^d \varphi^{(d)}$, so $\mathbb{E}_{s \sim N(0, \tau)} (\varphi \text{He}_d)(s) = (-1)^d (\varphi * \varphi_\tau)^{(d)}(0) = (-1)^d \varphi_{1+\tau}^{(d)}(0)$, and $\varphi_{1+\tau}(x) = r\varphi(rx)$ gives $\varphi_{1+\tau}^{(d)}(0) = r^{d+1} \varphi^{(d)}(0) = (-1)^d r^{d+1} (\varphi \text{He}_d)(0)$. Integrate against dr (Theorem 6.1). $\square$

Proposition 6.3 (Contact order is a concentration coefficient). *For a pre-activation with Gaussian surrogate $N(m, \sigma^2)$, $\alpha = m/\sigma$, the concentration function at the kink is*

$$Q(h) = \mathbb{P}(|z| \leq h) = 2f(0)h + \frac{1}{3}f''(0)h^3 + O(h^5),$$

$$f(0) = \frac{\varphi(\alpha)}{\sigma} = \frac{(\varphi \text{He}_0)(\alpha)}{\sigma}, \quad f''(0) = \frac{(\varphi \text{He}_2)(\alpha)}{\sigma^3}.$$

The coefficient of h^{d+1} is the order- d Hermite kink content. An estimator whose leading defect is an error in that coefficient has, by Theorem 6.2, merge constant $1/(d+2)$: the Euler–Stein constant reads off which Taylor coefficient of the concentration functions at the kinks the estimator gets wrong.

Proof. $Q(h) = \int_{-h}^h f = 2f(0)h + \frac{1}{3}f''(0)h^3 + O(h^5)$, and $\partial_z^k [\varphi((z-m)/\sigma)/\sigma]_{z=0} = (-1)^k \sigma^{-k-1} \varphi^{(k)}(-\alpha)$, with $\varphi^{(k)} = (-1)^k \text{He}_k \varphi$ and parity. $\square$

Conjecture 6.4 (Contact order at depth). *For the chains of stage 9 the defect at depth is dominated by one contact order, and $a = 1/(d+2)$ with: $d = 0$ for the Gaussian closure and for chains with truncated third-cumulant memory (their errors are variance errors acting through $f(0)$), $d = 2$ for the exact first-order third-cumulant chain (its errors are second-order, the curvature $f''(0)$).*

chain (official networks, 2049 exact chain runs each)	measured a	predicted	order d
Gaussian closure, network 0 / 1	0.535 / 0.500	1/2	0
window-one third-cumulant chain (memory dropped beyond one age)	0.482	1/2	0
exact first-order third-cumulant chain, network 0 / 1	0.255 / 0.257	1/4	2

The layer-wise profile of the absolute kink density along the localization was measured at all sixteen layers of network 0 to decay as $(1 + \tau)^{-1/2} = r$ (stage-9 workflow), which is the $d = 0$ profile; no profile of the exact chain's defect has been measured, so the $d = 2$ attribution rests on the constant. The earlier explanation of 0.255 as “two births, 1/3, reduced by rotation” is superseded: the order law gives 1/4 with no correction.

6.3 Mehler factors are imaginary modular times

Proposition 6.5 (Mehler–KMS). *Let N be the Hermite number operator on $L^2(\mathbb{R}, \gamma)$ ($N\text{He}_k = k\text{He}_k$), $\omega_\beta(Y) = \text{Tr}(\rho_\beta Y)$ with $\rho_\beta = (1 - e^{-\beta})e^{-\beta N}$ the Gibbs state of N , faithful and non-tracial. For a standard Gaussian pair (X, Y) with correlation $\rho = e^{-\beta} \in (0, 1)$ and real $f, g \in L^2(\gamma)$,*

$$\mathbb{E} f(X)g(Y) = \langle f, e^{-\beta N} g \rangle = \frac{1}{1 - \rho} \omega_\beta(|g\rangle\langle f|).$$

The modular group of ω_β is $\sigma_s(Y) = e^{-i\beta s N} Y e^{i\beta s N}$, the Mehler rotation (fractional Fourier transform of Gaussian space); the Ornstein–Uhlenbeck semigroup e^{-tN} is its continuation to imaginary time. The correlation of two Gaussian variables is thus the Boltzmann factor of one Hermite quantum at inverse temperature $\beta = -\log \rho$, and every pair correlation of the network, every Euler–Stein weight and every Wick coefficient of the chains is a matrix element of $e^{-\beta N}$ at some β .

Proof. Mehler’s formula: $\mathbb{E} f(X)g(Y) = \sum_k \rho^k \langle f, H_k \rangle \langle g, H_k \rangle$ with $H_k = \text{He}_k / \sqrt{k!}$; and $\text{Tr}(\rho_\beta |g\rangle\langle f|) = (1 - e^{-\beta}) \langle f, e^{-\beta N} g \rangle$. The modular group of a Gibbs state ρ_β is $\text{Ad } \rho_\beta^{is}$. $\square$

Proposition 6.6 (The ReLU heat trace and the folding exponent). *Let $K(\beta) = \mathbb{E}[\text{relu}(X) \text{relu}(Y)]$ at correlation $e^{-\beta}$, a weighted heat trace of the Hermite spectrum of relu : $K(\beta) = \sum_k c_k^2 e^{-\beta k} / k!$ with $c_0 = \frac{1}{\sqrt{2\pi}}$, $c_1 = \frac{1}{2}$, $c_{2j+2} = (-1)^j (2j - 1)!! / \sqrt{2\pi}$, $c_{\text{odd} \geq 3} = 0$. The weights decay like $k^{-5/2}$ (the kink), and*

$$K(\beta) = \frac{1}{2} \left(1 - \beta + \frac{2\sqrt{2}}{3\pi} \beta^{3/2} + O(\beta^2) \right), \quad \beta_{l+1} = \beta_l - \frac{2\sqrt{2}}{3\pi} \beta_l^{3/2} + O(\beta_l^2)$$

for the pair temperature $\beta_l = -\log \cos \theta_l$ of two inputs. The linear coefficient is exactly the critical one (the depth flow is marginal at linear order because relu is homogeneous: $K'(0)/K(0) = -\mathbb{E} \text{relu}^2 / \mathbb{E} \text{relu}^2 = -1$), and the kink’s $\beta^{3/2}$ term drives $\beta_l \simeq 9\pi^2 / (2l^2)$, i.e. $\theta_l \simeq 3\pi/l$. Depth heats the neuron pairs quadratically, $T_l = 1/\beta_l \simeq 2l^2 / (9\pi^2)$.

Proof. $\mathbb{E}[\text{relu}(Z)\text{He}_k(Z)] = \mathbb{E}[\text{relu}^{(k)}(Z)] = \mathbb{E}[\delta^{(k-2)}(Z)] = \text{He}_{k-2}(0)\varphi(0)$ for $k \geq 2$, and $\text{He}_{2j}(0) = (-1)^j(2j-1)!!$; Stirling gives $c_{2j+2}^2/(2j+2)! \sim (8\pi^{3/2})^{-1}j^{-5/2}$. In closed form $K = \frac{1}{2\pi}(\sin\theta + (\pi - \theta)\cos\theta)$ with $\cos\theta = e^{-\beta}$, whose non-analytic part $\frac{1}{2\pi}(\sin\theta - \theta\cos\theta) = \frac{\theta^3}{6\pi} + O(\theta^5)$ with $\theta = \sqrt{2\beta}(1 + O(\beta))$ is $\frac{\sqrt{2}}{3\pi}\beta^{3/2}$. The recursion is $e^{-\beta_{l+1}} = K(\beta_l)/K(0)$, and $\beta^{-1/2}$ grows by $\frac{\sqrt{2}}{3\pi}$ per layer. $\square$

Remark 6.7 (One clock). Three flows are one structure. Depth raises the pair temperature, $\beta_l \rightarrow 0$ like l^{-2} , with exponent fixed by the Weyl-type tail $k^{-5/2}$ of the activation's Hermite spectrum (a smooth activation has an analytic heat trace and flows exponentially unless tuned). Localization lowers it: the ball radius is the Boltzmann factor $r = e^{-t}$ of the Ornstein–Uhlenbeck time, and the Euler–Stein constant $1/(d+2) = \mathbb{E}e^{-(d+1)t}$ is the mean Boltzmann factor of $d+1$ quanta at an exponential time. The quantum Hermite transform of stage 9's reading is the real-time version of the same group. In this sense the Gaussian input's modular flow, the network's depth flow and the estimator's error clock are the imaginary-time, discrete-time and averaged-time faces of the single semigroup e^{-tN} .

7 Bootstrap: the mean as the unique solution of a positivity problem

Khalkhali and Pagliaroli [6] solve Dirac ensembles without solving them: Schwinger–Dyson (loop) equations give linear relations among moments, positivity of moment matrices cuts out a convex island, and the island shrinks onto the answer as the moment degree grows. The network admits the same treatment, with the Gaussian input as the “action”, Stein's identity as the loop equation, and the gates as idempotent variables.

Definition 7.1 (The ReLU moment problem). Let $\mathbb{R}[x, \theta]$ be the polynomial algebra in $x \in \mathbb{R}^n$ and commuting idempotents $\theta_{l,a}$ ($\theta^2 = \theta$), $l = 1, \dots, L$. Define recursively the polynomials $z_1 = W_1x$, $z_{l,a} = \sum_b W_l(a, b)\theta_{l-1,b}z_{l-1,b}$. A *feasible moment functional* is a linear $\mathcal{L} : \mathbb{R}[x, \theta] \rightarrow \mathbb{R}$ with $\mathcal{L}(1) = 1$, Gaussian x -moments (equivalently the loop equations $\mathcal{L}(x_i q) = \mathcal{L}(\partial_i q)$ for $q \in \mathbb{R}[x]$), positive moment matrix $\mathcal{L}(p^2) \geq 0$, and positive localizing forms $\mathcal{L}(p^2\theta_{l,a}z_{l,a}) \geq 0$, $\mathcal{L}(-p^2(1 - \theta_{l,a})z_{l,a}) \geq 0$ for all p .

Theorem 7.2 (The island is a point). *For almost every weight realisation every feasible moment functional represented by a measure is the law of $(x, \theta(x))$ except for the values of the gates on dead tiles, where all downstream pre-activations vanish and the gates are unconstrained; these measures are moment-determinate, and they all give the same value to every pre-activation moment. In particular $u = \mathcal{L}(\theta_L \circ z_L)$ is determined by positivity and the loop equations alone.*

Proof. A representing measure μ lives on $\mathbb{R}^n \times \{0, 1\}^{nL}$ and is supported in $K = \{\theta_{l,a}z_{l,a} \geq 0, (1 - \theta_{l,a})z_{l,a} \leq 0\}$. If $z_{1,a}(x) > 0$ the second constraint forces $\theta_{1,a} = 1$, if $z_{1,a}(x) < 0$ the first forces $\theta_{1,a} = 0$; given the gates of layers $< l$, $z_l(x, \theta)$ is the true pre-activation and the same argument forces layer l . So the fibre of K over x is the single point $\theta(x)$ unless some $z_{l,a}(x) = 0$. On a live tile this happens on a γ -null set ($z_{l,a}$ is a nonzero linear form there almost surely); on a dead tile, where $h_l \equiv 0$, all later pre-activations vanish whatever the later gates are, so the fibre is a set of gate patterns on which every $z_{l'}, l' > l$, and h_L take the value 0. Since the x -marginal is γ , μ is the graph measure up to the choice of these irrelevant gates. Determinacy: $\mu = \sum_{\vartheta} \mu_{\vartheta} \otimes \delta_{\vartheta}$ over the finitely many gate patterns, with $\mu_{\vartheta} \leq \gamma$; the moments of μ determine those of every μ_{ϑ} .

(the monomials in θ span the indicators of patterns), and each μ_θ satisfies Carleman's condition because its moments are bounded by the Gaussian ones. $\square$

Under the standard convergence theory of moment–sum-of-squares relaxations for determinate measures [16], the islands of degree $2k$ form a nested sequence of intervals containing u and shrinking to it. Two features are specific to the network.

Proposition 7.3 (Loop equations need the kink generators). *Stein's identity for test functions that involve gates produces the distributions $\delta(z_{l,a})$: $\partial_i \theta_{l,a} = \delta(z_{l,a}) \partial_i z_{l,a}$. The loop equations therefore close only on the algebra generated by x , the gates and the kink distributions $\delta^{(k)}(z_{l,a})$. The first of them, for the test field $q = J_i$, is the kink-current formula (1); the Wick coefficients $\mathbb{E}[\partial^k \text{relu}^p]$ of every cumulant chain are the Gaussian-state values of these kink generators.*

Proof. Differentiate $\theta_{l,a}(x) = \mathbf{1}[z_{l,a}(x) > 0]$ in the sense of distributions; the first loop equation is $\mathbb{E}[x \cdot \nabla h_L] = \mathbb{E}[\Delta h_L]$, computed in Theorem 4.1. For a Gaussian z , $\mathbb{E}[\delta^{(k)}(z)]$ are the derivatives of the density at the kink, which is what the chains' Wick tables contain. $\square$

Proposition 7.4 (The level-one island of a single neuron is Scarf's interval). *If only the mean m and variance s^2 of $z_{l,a}$ are imposed, together with $h \geq 0$, $h \geq z$, $h(h - z) = 0$, the closure of the set of feasible values of $\mathbb{E}h$ is exactly $[\max(0, m), \frac{1}{2}(m + \sqrt{m^2 + s^2})]$. The Gaussian closure picks the interior point $s(\alpha\Phi(\alpha) + \varphi(\alpha))$, $\alpha = m/s$.*

Proof. The constraints force $h = z_+$; the lower bound is Jensen, and the upper bound is Scarf's worst case over laws with given mean and variance, attained by a two-point law [15]. $\square$

So a moment chain is a truncated Schwinger–Dyson hierarchy solved with a Gaussian (or Edgeworth) ansatz and *without* positivity. Positivity is what would turn a chain into a certificate; its cost at width 1024 is prohibitive at any useful degree, and its role here is structural: it shows that the exact answer is the unique point of an island cut out by the input's loop equations and the complementarity of the gates, and that the information a chain discards is the positivity of the kink generators' moment matrices, which is exactly what goes wrong when a closure produces a negative fourth cumulant or a variance outside its admissible range.

8 What the placement explains, what it predicts, and what is open

8.1 What the placement explains

1. **Why the mean is hard and the densities are easy.** The mean is a sum over walls of a density at the kink times a conditional transport on the wall (Theorem 4.1). Densities are concentration-function coefficients and the Gaussian surrogate gets them right; the conditional transports are rare-event-relative objects (conditioning on a measure-zero wall), and every chain's error lives there.
2. **Why the third-cumulant memory has an identity leg.** It is the label of the kink at which the cumulant was born (Theorem 5.3); the memory is the record of all pinnings, transported.
3. **Why it cannot be compressed.** Kinks have small codegree at He initialisation, so their atoms are orthogonal; the frame bound (Theorem 5.4) makes the memory exactly as compressible as its transport leg, and the complementarity of the neuron algebra with the modular

frame of the layer state (Theorem 5.7) makes state-respecting (spectral) compressions blind to the hub diagonal up to $2\sqrt{2/n}$. The same small codegree that makes Sahasrabudhe’s nibble concentrate makes our compression fail.

4. **Why ensemble surrogates fail.** The memory’s readout is odd in the last weight matrix (Proposition 5.8); free and NNGP descriptions are the symmetric phase.
5. **Where the Euler–Stein constants come from.** The correction averages the defect over Gaussian balls of uniform radius (Theorem 6.1); a defect of Hermite contact order d decays like r^{d+1} , so $a = 1/(d+2)$, and d is the order of the concentration coefficient the estimator gets wrong (Theorem 6.2, Proposition 6.3); the three measured constants fit $d = 0, 0, 2$.
6. **Why depth behaves algebraically.** Lengths are preserved and chords contract like $3\pi/l$ (Theorem 2.4) because the activation’s heat trace is critical at linear order and has a $\beta^{3/2}$ kink term (Proposition 6.6).

8.2 What it predicts (falsifiable, theoretical)

1. The defect of the exact first-order chain along Eldan’s path decays like $r^3 = (1+\tau)^{-3/2}$ (order $d = 2$), not like $(1+\tau)^{-1}$; a chain that also corrects the curvature coefficient $f''(0)$ at the kinks has a defect of order $d \geq 4$ and merge constant at most $1/6$.
2. On the official networks, the number of walls crossed by a great-circle arc grows linearly in depth (about $n\theta_0/\pi$ per layer) while the number of gates that differ between its endpoints grows like $3n \log L$.
3. The best hub-index compression of any old source has the error of the low-rank truncation of its transport leg $T_{l \leftarrow \nu} C$ (Theorem 5.4 with a near-isotropic Hadamard leg); a tier cannot do better than that number, which can be computed without running any chain.
4. The complementarity constant of the layer state at depth is $\approx 2\sqrt{2/n}$ up to the collective mode, so every state-preserving coarse-graining of a layer onto spectral data retains at most a fraction $\approx 8/n$ of the centred energy of any neuron-diagonal observable (a gate pattern, a hub diagonal).

8.3 Open problems

1. **Contact order at depth.** Prove Conjecture 6.4: show that the defect of a chain whose first neglected term is of Hermite order d at every wall has the profile r^{d+1} along the localization, despite the nonlinear dependence of the downstream pre-activation means on the localization centre.
2. **A contact-resolved estimator.** The kink-current formula asks for the on-wall transports $\mathbb{E}[\partial_{h_l} h_L e_a \mid z_{l,a} = 0]$ for all walls. To first order these are pinning responses (Lemma 5.2); their sum over walls is one directional derivative of the downstream chain per output, which forward mode computes with an identity leg (the memory) and reverse mode with one adjoint per output. Either way the cost is that of the memory. Is there an order-two quantity (a pinning response of pinning responses) whose sum over walls collapses, as the radial fourth-cumulant scalar does for the fourth cumulant?

3. **The induced Dirac ensemble.** Compute the spectral dimension and spectral variance [7] of the selected Dirac operators $P_x D P_x$ and of the transports $T_{l \leftarrow l'}$ by free multiplicative convolution, and compare with the measured stable ranks 164, 92, 65, 47, 29, 23 of the cocycle at ages 1, 2, 3, 4, 6, 8; decide whether the emergence of the collective eigenvalue of the layer covariance (from 0.9% to 37–48% of the energy) is a Baik–Ben Arous–Péché transition in depth, the analogue of the one-cut to two-cut transition of [5].
4. **Bootstrap with kink generators.** Is the island of the moment problem at low degree narrow at large n because of a factorisation analogous to large- N factorisation [6]? If so, positivity would certify chains cheaply.
5. **Klartag in reverse.** Klartag’s ellipsoid stops when about $d^2/2$ contacts hold its $d(d+1)/2$ shape parameters. The marginal value of a source age decays by a factor 0.75 per age and about eleven ages matter, roughly $4n$ effective contacts held in memory. Is there a counting argument that fixes this number from the dimension of the third-cumulant readouts that the mean actually uses?

Verification

`scripts/verify_stage10.py` compares the two sides of each identity numerically; it runs no estimator and no network of the competition. Output in `notes/stage10/verify_stage10.txt`.

statement	check	result
Theorem 4.1 (kink current)	$\mathbb{E} h_L$ against the wall sum with a ± 0.05 window, $3 \cdot 10^6$ inputs, $L = 3$	$n = 16$: agreement to noise (1.5%); $n = 6$: 25% off, dead tiles carry 16%
Theorem 2.2 (Crofton)	crossings of $2 \cdot 10^4$ random walls by a random curve in $\mathbb{R}^{50}$	31.28 against $\Lambda/\pi = 31.39$
Theorem 2.4, Proposition 6.6	iterate the angle map	$\theta_l l / 3\pi = 0.90, 0.985, 0.992$ at $l = 10^2, 10^3, 2 \cdot 10^3$ (log corrections)
Proposition 2.3 (unit gain)	tangent and norm gains at $\sigma_w^2 = 2$, $n = 400$	1.009, 1.002
Proposition 6.5 (Mehler, Hermite coefficients of relu)	closed form against the Hermite series	equal to 10^{-8}
Theorem 6.2 (contact order)	stage-9 τ -integral for $d = 0, 2, 4$	0.497, 0.246, 0.164 against $\frac{1}{2}, \frac{1}{4}, \frac{1}{6}$ (truncated tail)
Theorem 5.3 (births are contacts)	κ_3/ε^2 by Sobol quasi-Monte Carlo, 2^{23} points	ratio to the hub formula 1.064, 1.030, 1.015, 1.009 at $\varepsilon = 0.16, \dots, 0.02$
Theorem 5.7 (complementarity)	$\ E_B E_A - J\ $ against $\ C\ $ for $S = WW^\top$	equal; 0.333 / 0.173 against $2\sqrt{2/n} = 0.354$ / 0.177 ($n = 64, 256$)
Theorem 5.4 (frame bound)	2000 random rank-10 compressions	bound holds

The asymptotic $3\pi/l$ is reached slowly ($O(\log l/l)$ corrections), so at the depth $L = 16$ of the competition the exact recursion, not its asymptotic form, is the statement to use.

References

- [1] Stage 9 note: *the mean of a random ReLU network as a central localization* (this repository, [notes/stage9](#)).
- [2] Stage 8 note: *Localization in the commutant: modular tilts, Kikuchi–Toeplitz quantization and transverse measures*.
- [3] J. Sahasrabudhe, Probabilistic combinatorics at exponentially small scales, [arXiv:2512.15077](#).
- [4] B. Klartag, Lattice packing of spheres in high dimensions using a stochastically evolving ellipsoid, [arXiv:2504.05042](#).
- [5] M. D’Arcangelo, S. Gnutzmann, Symmetry breaking and phase transitions in random non-commutative geometries and related random-matrix ensembles, [arXiv:2601.14141](#).
- [6] M. Khalkhali, N. Pagliaroli, Bootstrapping noncommutative geometry with Dirac ensembles, [arXiv:2512.08694](#).
- [7] J. W. Barrett, P. Druce, L. Glaser, Spectral estimators for finite non-commutative geometries, [arXiv:1902.03590](#).
- [8] J. W. Barrett, L. Glaser, Monte Carlo simulations of random non-commutative geometries, *J. Phys. A* 49 (2016) 245001.
- [9] R. Wang, L. C. Lau, H. Zhou, Derandomizing matrix concentration inequalities from free probability, [arXiv:2601.08111](#).
- [10] A. Sah, J. Sahasrabudhe, M. Sawhney, On the Spielman–Teng conjecture, [arXiv:2405.20308](#).
- [11] Y. Cho, L. K. Saul, Kernel methods for deep learning, *NeurIPS* 2009.
- [12] B. Poole, S. Lahiri, M. Raghu, J. Sohl-Dickstein, S. Ganguli, Exponential expressivity in deep neural networks through transient chaos, *NeurIPS* 2016.
- [13] S. Hayou, A. Doucet, J. Rousseau, On the impact of the activation function on deep neural networks training, *ICML* 2019.
- [14] B. Hanin, D. Rolnick, Complexity of linear regions in deep networks, *ICML* 2019.
- [15] H. Scarf, A min-max solution of an inventory problem, in *Studies in the Mathematical Theory of Inventory and Production*, Stanford 1958.
- [16] J.-B. Lasserre, Global optimization with polynomials and the problem of moments, *SIAM J. Optim.* 11 (2001) 796–817.
- [17] M. Takesaki, Conditional expectations in von Neumann algebras, *J. Funct. Anal.* 9 (1972) 306–321.
- [18] M. A. Rieffel, Metrics on state spaces, *Doc. Math.* 4 (1999) 559–600.
- [19] A. Connes, *Noncommutative Geometry*, Academic Press 1994.